

Binary Search Trees

Instructor: Meng-Fen Chiang

COMPCSI220: WEEK 9



Slides adapted from Mark Wilson, Georgy Gimel'farb, Simone Linz and Tanya Gvozdeva

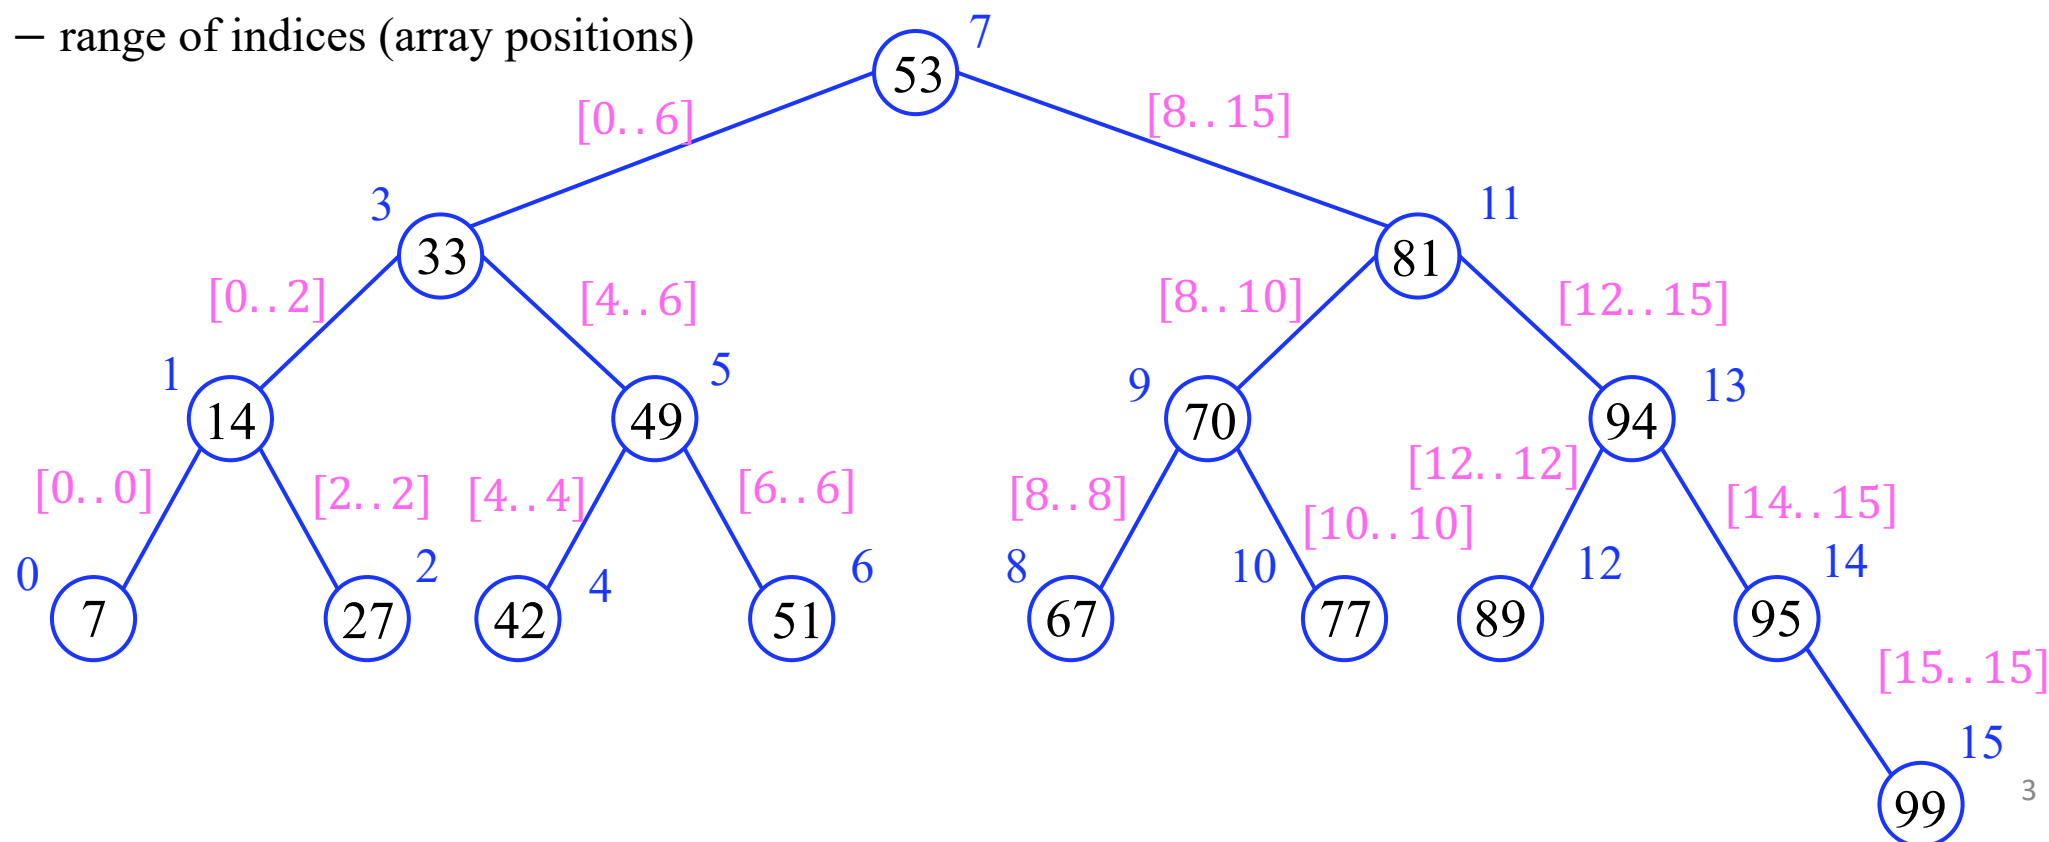
OUTLINE

- Tree Data Structure
- Binary Search Tree Operations
- Time Complexity Analysis

Tree Structure of Binary Search

| | | | | | | | | | | | | | | | |
|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 |
| 7 | 14 | 27 | 33 | 42 | 49 | 51 | 53 | 67 | 70 | 77 | 81 | 89 | 94 | 95 | 99 |



$[l..r]$ – range of indices (array positions)




Binary Search Tree: Left-Right Ordering of Keys

- Left-to-right numerical ordering in a BST: for every node i ,
 - the values of all the keys $k_{\text{left}:i}$ in the left subtree are smaller than the key k_i in i and
 - the values of all the keys $k_{\text{right}:i}$ in the right subtree are larger than the key k_i in i :


$$\{k_{\text{left}:i}\} \ni l < k_i < r \in \{k_{\text{right}:i}\}$$



Binary Search Tree: Left-Right Ordering of Keys




BST



Non-BST:

Key "2" cannot be in the right subtree of key "3".



Non-BST:

Key "11" and "12" cannot be in the left subtree of key "10".

Binary Search Tree Operations


- BST is an explicit data structure implementing the table ADT.
 - BST are more complex than heaps: any node may be removed, not only a root or leaves.
 - The only practical constraint: no duplicate keys (attach them all to a single node).
- Basic operations
 - **Find** a given search key or detect that it is absent in the BST.
 - **Insert** a node with a given key to the BST if it is not found.
 - **FindMin**: find the minimum key.
 - **findMax**: find the maximum key.
 - **Remove** a node with a given key and restore the BST if necessary.

BST Operations: Find / Insert a Node

find: a successful binary search


insert: creating a new node at the point where an unsuccessful search stops.

Example: Find 8



Found node

Example: Insert 7



Inserted node

BST Operations: FindMin / FindMax

- Extremely simple: starting at the root, branch repeatedly left (**findMin**) or right (**findMax**) as long as a corresponding child exists.
- The **root of the tree** plays a role of the **pivot** in quicksort and quickselect.
- As in quicksort, the **in-order traversal** of the tree can sort the items:
 - First visit the left subtree;
 - Then visit the root, and
 - Then visit the right subtree.
- $O(\log n)$ average-case and $O(n)$ worst-case running time for **find**, **insert**, **findMin**, and **findMax** operations, as well as for **selecting** a single item

BST Operations: Remove a Node

- The most complex because the tree may be disconnected. Need to reconnect some nodes!
 - Reconnection must retain the ordering condition.
 - Reconnection should not needlessly increase the tree height.

BST Operations: Remove a Node

- Standard method of removing a node i with c children:


| c | ACTION |
|---|--|
| 0 | Simply remove the leaf i . |
| 1 | Remove the node i after linking its child to its parent node. |
| 2 | Swap the node i with the node j having the smallest key k_j in the right subtree of the node i . After swapping, remove the node i (as now it has at most one right child). |

BST Operation: Remove a Node

Remove 3



Link 1 to the left branch of 5






BST Operation: Remove a Node

Remove 10



Replace 10 (swap with 12 and delete)




Minimum key in the right subtree

The Worst-Case Time Complexity


- The find, insert, and remove operations in a BST all take time in $O(h)$ in the **worst case**, where h is the height of the tree.
- **Proof:** The running time $T(n)$ of these operations is proportional to the number of nodes visited.
 - **Find / insert:** $1+h$.
 - **Remove:** "1 + the depth of the node + the height of its highest subtree" $\rightarrow 1+h$.
 - In each case $T(n) = \Theta(h)$.
 - For a well-balanced BST, $T(n) \in O(\log n)$ (logarithmic time).
 - In the worst case $T(n) \in \Theta(n)$ (linear time) because insertions and deletions may heavily destroy the balance.

The Worst-Case Time Complexity

BSTs of height $h \approx \log n$




BSTs of height $h \approx n$



The Average-Case Time Complexity


- More balanced trees are more frequent than unbalanced ones.
- **Definition (Internal Path Length):** The total internal path length, $S_\tau(n)$, of a binary tree τ is the sum of the depths of all its nodes.



- Average complexity of a successful search in τ : the average node depth, $1/n S_\tau(n)$, e.g. $1/8 S_\tau(8) = 15/8 = 1.875$ in this example.

The Average-Case Time Complexity (Contd.)

- Average-case complexity of searching:
 - Averaging $S_{\tau}(n)$ for all the trees of size n , i.e. for all possible $n!$ Insertion orders, occurring with equal probability, $\frac{1}{n!}$



$$\begin{aligned}
 S_{\tau}(8) &= 0 + \\
 &\quad + 1 + 1 \\
 &\quad + 2 + 2 \\
 &\quad + 3 + 3 + 3 = \mathbf{15}
 \end{aligned}$$

The $\Theta(\log n)$ Average-case BST Operations

- Let $S(n)$ be the average of the total internal path length, $S_\tau(n)$, over all BST τ created from an empty tree by sequences of n random insertions, each sequence considered as equally possible.
- The expected time for successful and unsuccessful search (insertion and deletion) in such BST is $\Theta(\log n)$
- **Proof**: It should be proven that $S(n) \in \Theta(n \log n)$
 - Obviously, $S(1) = 0$.
 - Any n -node tree, $n > 1$, contains a left subtree with i nodes, a root at level 0, and a right subtree with $n - i - 1$ nodes; $0 \leq i \leq n - 1$.
 - For a fixed i , $S(n) = (n - 1) + S(i) + S(n - i - 1)$, as the root adds 1 to the path length of each other node.

The $\Theta(\log n)$ Average-case BST Operations (Contd.)

- After summing these recurrences for $0 \leq i \leq n - 1$ and averaging, just the same recurrence as for the average-case quicksort analysis is obtained:

$$S(n) = (n - 1) + \frac{2}{n} \sum_{i=0}^{n-1} S(i)$$

- Therefore, $S(n) \in \Theta(n \log n)$, and the expected depth of a node is $\frac{1}{n} S(n) \in \Theta(\log n)$.
- Thus, the average-case search and insertion time is in $\Theta(\log n)$.
- It is possible to prove (but in a more complicated way) that the average-case deletion time is also in $\Theta(\log n)$.
- The BST allow for a special **balancing**, which prevents the tree height from growing too much, i.e. avoids the worst cases with linear time complexity $\Theta(n)$.

SUMMARY

- Tree Data Structure
- Binary Search Tree Operations
 - find, insert, and remove operations
- Time Complexity Analysis